

Towards a Context Layer for Self-Improving Data Agents

Till Döhmen
MotherDuck
Netherlands
till@motherduck.com

Jacob Matson
MotherDuck
USA
jacob@motherduck.com

Jordan Tigani
MotherDuck
USA
jordan@motherduck.com

Abstract

Data agents that autonomously query and analyze structured data struggle with accuracy when they lack organizational context: which table is canonical, how metrics are defined, which join patterns are correct. We present results from a benchmarking study on the DABstep text-to-SQL benchmark, where providing domain documentation and targeted data modeling (views, macros) improved hard-question accuracy from 29.8% to 93.2% while reducing cost by 62%. Building on these findings, we describe a *context layer*: a data-platform-native knowledge store where agents both consume and contribute organizational knowledge, and which serves as the foundation for targeted data engineering that we envision can eventually be largely automated. We argue that the key challenges are bridging the gap between personal and organizational memory, building effective agentic retrieval beyond naive RAG, and enabling organizations to evaluate their agents against their own data through self-serving evals. As models converge in capability, the differentiator for data agents will increasingly be the context they have accumulated, not which model they call.

Keywords

data agents, context engineering, text-to-SQL, agentic memory, retrieval, self-improvement

1 Introduction

Today’s data systems were designed for a small number of careful human operators writing hand-crafted SQL. A growing share of analytical workloads is now delegated to AI agents [13] that translate natural-language questions into SQL, build visualizations, and orchestrate data pipelines. This shift exposes a critical gap: agents lack the organizational context that human analysts accumulate over years. Which table is the source of truth? How is “revenue” computed? Which filters are always required? Yet even providing context is only effective to a certain degree; the last mile of accuracy requires targeted data modeling that encodes complex logic directly in the database.

The industry has recognized this context gap. a16z calls context “the key missing ingredient” for data agents [11]. OpenAI built a six-layer context system for their in-house data agent [18]. GitHub Copilot shipped cross-agent memory [9]. Yet these solutions are either proprietary, generic (unaware of data models and organizational permissions), or both.

This paper makes three contributions. First, we provide background on context for data agents, from semantic layers to agentic memory, and argue that the problem is not adequately served by existing approaches (§2). Second, we present benchmarking results showing that domain documentation and targeted data modeling have more impact on agent accuracy than the choice of model (§3).

Third, we describe a possible design of a *context layer* and identify the key open problems: bridging personal to organizational memory in a multi-user setting, building effective agentic retrieval, and enabling domain-specific evaluation (§4). Finally, we outline a vision for the fully agentic data stack (§5).

2 Background

2.1 Semantic Layers

Semantic layers act largely as an intermediate translation layer that guarantees semantic correctness of queries: business metrics, dimensions, and relationships are defined in a structured markup language, and a query engine translates queries formulated in the semantic layer’s own query language into SQL (e.g. dbt Semantic Layer [6], Cube [5], Malloy [14]). Implementations vary and there is a push toward standardization [17], but the basic principle remains the same.

In practice, many questions in an agentic text-to-SQL context are exploratory and go beyond pre-defined metrics. Business definitions evolve continuously, and maintaining a semantic layer that reflects reality is cumbersome, as we have argued previously [16]. We argue that a better division of labor is: things that are known and stable should be modeled on the database layer itself (as views, macros, and comments that agents discover through standard schema exploration), while everything beyond should be captured in a self-improving memory system that we call a context layer, optimized for continuous learning and agentic retrieval.

2.2 Agentic Memory

The AI community has converged on the idea that agents need persistent memory, and the field has matured rapidly from 2024–2026. **General-purpose memory systems** like Mem0 [4], Hind-sight [12], and Cognee [15] are converging on similar architectures: hybrid storage (vector + graph), multi-strategy retrieval with reranking, explicit temporal modeling, and active memory consolidation. Multi-strategy retrieval (semantic, keyword, graph traversal, temporal) is the single biggest performance differentiator, and temporal reasoning remains the hardest unsolved capability. General-purpose memory, however, does not adequately address the needs of data agents. For example, knowledge lifecycles are domain-specific: in code, a memory should expire when the cited file is deleted or move when it is renamed; in data, knowledge must stay in step-lock with schema evolution and decay when tables are deprecated. Access control is also domain-specific: code memory should mirror repository permissions, data memory should mirror data warehouse permissions. Generic memory systems have no mechanism for either.

We see domain-specific memory systems emerging in the SWE domain, like GitHub Copilot’s cross-agent memory [9], which stores

observations with citations to code locations and verifies them in real-time against the current branch, tying knowledge decay to the code objects themselves. And we see organization-specific solutions for data agent memory emerging across the industry. One widely covered example is OpenAI’s data context layer [18], which combines table lineage, human annotations, code-level enrichment from pipeline source, institutional documents, a personal/global memory layer, and live runtime queries, all tightly coupled to their 70k-table warehouse and permissions model.

We see an open need for a memory architecture that is specific enough to address the domain needs of data agents (*object-scoped*, tied to databases, tables, and columns; *organizationally permissioned*; *trust-differentiated*, where an admin-endorsed definition outranks a personal note), yet generic enough to not be tied to a single organization’s stack.

3 Benchmarking Context Engineering

We conducted a simple experiment, using a challenging text-to-SQL benchmark dataset, to motivate the need for both accurate domain context as well as targeted data modeling.

3.1 Experimental Setup

For the experiment, we used DABstep [7], a suite of 450 payment-processing questions¹ that comes with domain documentation and requires understanding of a complex fee-matching system with 9 join dimensions, wildcard joins, and hypothetical simulations. The original paper reports that “even the best agent achieves only 14.55% accuracy on the hardest tasks” with default setups.

All configurations use a benchmark-agnostic agent framework equipped with 3 MCP [1] tools (SQL execution, list tables, list columns) and a turn limit of 15 tool calls. We used a low-cost state-of-the-art LLM, Gemini 3 Flash [10], and varied two independent dimensions:

- (1) **Context delivery:** How domain knowledge reaches the agent. We tested three approaches:
 - (a) *Annotations:* hand-crafted column comments plus auto-generated statistical profiles, discoverable via the list-columns tool.
 - (b) *BM25 retrieval:* BM25 search over chunked domain documentation and SQL pattern examples, injecting relevant fragments into the prompt at query time.
 - (c) *Full-context prompt:* all domain documentation, term mappings, and SQL patterns injected into the system prompt in full.
- (2) **Data modeling:** Whether complex join logic is pre-computed as database objects. Views encode the 9-dimension fee-matching join; DuckDB table macros encapsulate parameterized calculations as callable functions. These were developed iteratively through qualitative evaluation of benchmark results, with a human and agent jointly identifying the highest-leverage views and macros to add.

¹DABstep contains 460 questions total, but we excluded the 10 development questions that ship with gold answers, as they are likely present in model training data.

3.2 Results

Table 1 shows accuracy across combinations of the two dimensions. The best configuration (full-context prompt + views/macros) achieves 93.2% on hard questions, which as of March 2026 would rank first on the DABstep leaderboard.²

Table 1: Accuracy on DABstep by context retrieval and data modeling (Gemini 3 Flash).

Context retrieval	Data modeling	Easy	Hard	Cost
None (baseline)	×	79.2%	29.8%	\$11.35
Column annotations	×	88.9%	30.1%	\$10.43
Full-context prompt	×	90.3%	82.7%	\$6.39
Full-context prompt	✓	93.1%	93.2%	\$4.35
BM25 retrieval	✓	91.7%	86.6%	\$5.96

3.3 Key Findings

The first finding concerns context delivery. Column annotations alone are insufficient: per-column comments added only 0.3% on hard questions because they cannot express cross-table logic. Retrieving chunks of domain documentation via BM25 performs better but remains lossy: keyword matching sometimes filters out patterns the model would have found with full visibility, and BM25 retrieval (86.6%) trails the full-context prompt (93.2%) by 6.6 points while costing 27% more (due to a higher number of turns). The full-context prompt is effective because it lets the model make its own relevance judgments, but dumping all documentation into a single prompt is only feasible because this benchmark has a bounded domain. For real-world organizations with thousands of tables, more sophisticated retrieval methods become necessary.

The second finding is that *data modeling* makes a significant contribution to the results. The fundamental bottleneck is reconstruction: LLMs can read business rules in prose but cannot reliably assemble them into complex multi-table joins under turn pressure. Views and macros solve this by moving logic to DDL time (+10.5 points when adding macros to the full-context prompt). Without data modeling, prompt engineering hits a ceiling at 82.7% on hard questions regardless of how well instructions are written. An open question is how to arrive at signals that inform data modeling decisions. This motivates a notion of trust and canonicalness in the context layer, which can form the basis for automated data modeling.

4 A Context Layer for Data Agents

Our benchmark shows that both rich context and targeted data modeling significantly improve agent accuracy. The question is how to establish context, make it retrievable, and do so in a way that can also inform data modeling decisions. An important starting point is to collect signals about how data is used and in which context it is used across the organization. In many cases, such

²Our scoring uses a custom evaluator derived from inspecting accepted Hugging Face submissions: bidirectional prefix matching for strings, strict float comparison, sign-insensitive numerics, and quote stripping. We have not yet submitted to the official DABstep evaluator; a submission is planned.

knowledge is not explicitly written down and does not fit in a single prompt at scale. Inspired by the agentic memory systems surveyed in §2, we propose one possible design for a self-maintaining *context layer* that reframes context as a *discovery problem* rather than a *definition problem*. We suggest such a system requires the following characteristics:

- **Object-scoped and multi-tenant:** Context must be tied to data objects (tables, columns, schemas) and scoped to users, roles, or the organization, inheriting the database’s access control model. The system must distinguish personal observations from organizational facts, with mechanisms to promote context across visibility levels.
- **Self-curating:** The system must consolidate, deduplicate, and validate raw context, accumulating from agent interactions and validating through evaluation to prevent noise and improve over time.
- **Agentically retrievable:** Delivery must go beyond pure search, supporting agentic retrieval methods like progressive disclosure and task-aware retrieval.

4.1 Fragments

The core unit of our envisioned context layer is the **fragment**: a plain-text piece of knowledge with metadata including structured references to database objects (databases, schemas, tables, columns) and other artifacts. Fragments can represent anything from a formal metric definition to an informal observation about data quality. What distinguishes them from unstructured documents is that they are *object-scoped* and *access-controlled*: each fragment is anchored to a specific data asset or set of data assets (tables, columns, etc.) and scoped to a user, role, or organization. This enables both better surfacing during retrieval and easier maintenance as schemas evolve.

Fragments are created through multiple channels (Figure 1). They can be *explicitly authored* by humans (e.g. metric definitions, business rules). *Ambient capture* records observations from agent conversations: every time an agent discovers a join pattern, corrects a misunderstanding, or receives feedback, the observation is stored as a fragment. *Inferred context* is derived directly from the data platform: query history mining [8], statistical profiling, and schema change detection produce knowledge that no human would think to write. *External sources* such as dbt model documentation, Slack threads, and Notion pages contain organizational knowledge that already exists but is scattered and inaccessible to agents. In all cases, fragment creation is typically LLM-mediated: a synthesis step extracts structured knowledge from raw material and anchors it to the relevant database objects.

4.2 Self-Curation

All input channels produce raw material that must pass through *curation and integration* before becoming trusted context (Figure 1). We envision a trust lifecycle: raw observations are unverified (rumors), consolidated and deduplicated context has been validated (lore), and human-endorsed definitions carry organizational authority (canon). Without this distinction, the context layer would fill with noise and contradictions. A *promotion* mechanism moves context from personal to organizational scope; sensitive personal




Figure 1: Architecture of the context layer. Raw observations from agent conversations (ambient capture), inferred context from the data platform (query history, profiling), and external sources (dbt models, documentation tools) all flow through a curation and integration step before entering the context layer. Human governance and evaluation form a bidirectional loop that promotes, endorses, and validates knowledge.

memory must not leak without deliberate action. The approach may combine LLM-driven consolidation [2] with human review.

For example, if three analysts independently observe that a column should never be summed directly because rows duplicate per fee rule, the system should consolidate these into a single endorsed fragment anchored to that column, rather than retaining three redundant personal observations.

4.3 Agentic Retrieval

Naive RAG (embed documents, similarity search, stuff into prompt) falls short for structured data. Relevance depends on more than textual similarity: object-level references, trust status, recency, and the agent’s current task all matter. This suggests that retrieval itself should be agentic: a retrieval agent that reasons about what context would be most useful for the current task [3].

Because fragments are object-scoped, retrieval can leverage the data model directly: when an agent queries a specific table, all fragments anchored to that table and its columns are surfaced automatically without embedding search. As the agent narrows its focus (e.g. from exploring a schema to writing a specific join), retrieval progressively discloses more specific fragments. This is fundamentally different from document retrieval, where relevance must be inferred from textual similarity alone.

4.4 Evaluation and Data Modeling

As the context layer grows, human oversight at the fragment level becomes impractical: there are too many fragments, they evolve too quickly, and judging individual fragment quality in isolation is difficult. We suggest leaving curation largely to agents, and providing oversight and observability through organization-specific **evaluations**.

Public benchmarks like DABstep measure agent capability in general, but cannot tell a user whether *their* agent understands *their* data. We call these **self-serving evals**: question-answer pairs

written against a user’s actual database, testing whether the agent returns correct results given currently available context. Instead of curating individual fragments, the user maintains a set of questions that are evaluated for correctness as the context layer evolves, giving visibility into how good the context layer is. An agent could also suggest question-answer pairs based on usage patterns in the organization. Self-serving evals serve as both accuracy measurements and regression tests when context is promoted from personal to organizational scope.

Beyond evaluation, fragment usage statistics (which fragments are retrieved, how often, for which query patterns) provide another signal for informing data modeling decisions. Frequently used fragments that encode complex join logic are natural candidates for, e.g., materialization as views or macros. This closes the loop between context and the data layer.

5 Outlook: The Agentic Data Stack

5.1 The Modern Data Stack Collapses

LLMs are already proficient at writing SQL, and on our platform, agent-generated queries now exceed human-authored ad-hoc queries. When humans stop writing SQL, much of the traditional data stack loses its *raison d’être*: BI tools, transformation layers, and orchestration frameworks were built for human operators. AI blurs the swim-lanes between these categories, and anything that can be vibe-coded will be [19]. What remains is compute (analytics is resource-intensive and agents will call out to query engines), and *context*: the organization-specific map between raw data and the questions an agent can answer reliably. Context cannot be inferred by the model alone; it must be accumulated, maintained, and delivered by the platform.

5.2 Stages of Agentic Analytics

We outline a progression through increasingly autonomous stages (described more deeply in [20]), illustrated in Figure 2:

- (1) **AI-assisted SQL**: Autocomplete and syntax help for human-written queries.
- (2) **Agentic text-to-results**: The agent runs multiple queries to answer a natural-language question, discovering schema as it goes.
- (3) **Agentic BI with context**: Agents produce queries and visualizations, guided by a curated context layer that ensures correct calculations.
- (4) **Self-driving context**: The context layer is maintained automatically. Agents capture knowledge ambiently, consolidation processes curate it, and humans review and endorse.
- (5) **Self-driving data models**: Accumulated context informs automated data modeling. Agents propose views, macros, and materializations that encode recurring join patterns and business logic permanently in the database layer.
- (6) **Autonomous insights**: Agents monitor data proactively, detect anomalies, generate dashboards, and surface hypotheses about root causes, all before a human asks.

The whole point of an analytics system is to provide answers to questions. Both context and the data model facilitate how well




Figure 2: Stages of agentic analytics [20]. The dashed line marks roughly where the industry stands today. The context layer (§4) is our proposal for enabling stages 4 and 5.

agents can do this. We believe stages 1–3 are well-understood today. The context layer described in §4 is our proposal for enabling stage 4. Stage 5 closes the loop: a mature context layer that tracks which patterns recur can inform data modeling decisions automatically, and automated data modeling in turn makes onboarding new data sources easier. At stages 5–6, specialized agents coordinate data quality monitoring, data model evolution, and context curation; we outline one such coordination model in [20]. While speculative, one plausible trajectory is that data engineers become curators of context and data quality, managing change and writing evals that serve as unit tests for the data. The analyst role blends with data engineering at higher stages of autonomy, shifting from writing queries to governing what questions the system should be answering. The final stage would invert the relationship entirely: the system already knows what questions matter and proactively provides answers before a human asks.

6 Conclusion

Our benchmarks show that both domain context and targeted data modeling significantly improve agent accuracy, even with a low-cost model. Building on this finding, we propose a context layer that makes organizational knowledge a first-class, self-maintaining resource. Open problems that we are actively exploring include bridging personal and organizational memory, building retrieval that goes beyond naive search, and giving organizations evaluation tools to measure whether their agent actually understands their data. We believe that as the primary consumers of data systems shift from humans to agents, agentic maintenance of context and data models becomes increasingly important, and the human skill required will eventually shift from “can write SQL” to “knowing what the system should care about.”

References

- [1] Anthropic. 2024. Model Context Protocol. <https://modelcontextprotocol.io/>
- [2] Anthropic. 2025. Claude Code Memory and Dream Consolidation. <https://milvus.io/blog/claude-code-memory-memsearch.md> Inspired by the Dream Consolidation Pattern.
- [3] Ash Ashutosh and Edo Liberty. 2026. Pinecone Nexus: The Knowledge Engine for Agents. <https://www.pinecone.io/blog/knowledge-infrastructure-for-agents/> Pinecone Blog.
- [4] Prateek Chhikara et al. 2025. Mem0: Building Production-Ready AI Agents with Scalable Long-Term Memory. *arXiv preprint arXiv:2504.19413* (2025). <https://arxiv.org/abs/2504.19413>

- [//arxiv.org/abs/2504.19413](https://arxiv.org/abs/2504.19413)
- [5] Cube Dev. 2026. Cube. <https://cube.dev/>
 - [6] dbt Labs. 2026. dbt Semantic Layer. <https://www.getdbt.com/>
 - [7] Alex Egg, Martin Iglesias Goyanes, Friso Kingma, Andreu Mora, Leandro von Werra, and Thomas Wolf. 2025. DABstep: Data Agent Benchmark for Multi-step Reasoning. *arXiv preprint arXiv:2506.23719* (2025). <https://arxiv.org/abs/2506.23719>
 - [8] Evgeniia Egorova. 2025. *Automatic Metadata Extraction for Text-to-SQL*. Master's thesis. <https://arxiv.org/abs/2505.19988> Master's thesis, in collaboration with MotherDuck.
 - [9] GitHub. 2025. Building an Agentic Memory System for GitHub Copilot. <https://github.blog/ai-and-ml/github-copilot/building-an-agentic-memory-system-for-github-copilot/>
 - [10] Google. 2026. Gemini 3 Flash. <https://deepmind.google/technologies/gemini/>
 - [11] Andreessen Horowitz. 2025. Your Data Agents Need Context. <https://a16z.com/your-data-agents-need-context/>
 - [12] Chris Latimer, Nicolás Boschi, Andrew Neeser, Chris Bartholomew, Gaurav Srivastava, Xuan Wang, and Naren Ramakrishnan. 2025. Hindsight is 20/20: Building Agent Memory that Retains, Recalls, and Reflects. *arXiv preprint arXiv:2512.12818* (2025). <https://arxiv.org/abs/2512.12818>
 - [13] Shu Liu, Soujanya Ponnappalli, Shreya Shankar, Sepanta Zeighami, Alan Zhu, Shubham Agarwal, Ruiqi Chen, Samion Suwito, Shuo Yuan, Ion Stoica, Matei Zaharia, Alvin Cheung, Natacha Crooks, Joseph E. Gonzalez, and Aditya G. Parameswaran. 2026. Supporting Our AI Overlords: Redesigning Data Systems to be Agent-First. In *CIDR*. <https://arxiv.org/abs/2509.00997>
 - [14] Malloy. 2026. Malloy. <https://www.malloydata.dev/>
 - [15] Vasilije Marković, Lazar Obradović, Laszlo Hajdu, and Jovan Pavlović. 2025. Optimizing the Interface Between Knowledge Graphs and LLMs for Complex Reasoning. *arXiv preprint arXiv:2505.24478* (2025). <https://arxiv.org/abs/2505.24478>
 - [16] Jacob Matson. 2025. What If We Don't Need the Semantic Layer? *MotherDuck Blog* (2025). <https://motherduck.com/blog/what-if-we-dont-need-the-semantic-layer/>
 - [17] Open Semantic Interchange. 2025. OSI Core Metadata Specification. <https://github.com/open-semantic-interchange/OSI> Version 0.1.1.
 - [18] OpenAI. 2025. Inside Our In-House Data Agent. <https://openai.com/index/inside-our-in-house-data-agent/>
 - [19] Jordan Tigani. 2026. Future Casting the Modern Data Stack. *MotherDuck Blog* (2026). <https://motherduck.com/blog/future-casting-the-modern-data-stack/>
 - [20] Jordan Tigani. 2026. Water-Town: The Agent Swarm Data Stack. *MotherDuck Blog* (2026). <https://motherduck.com/blog/water-town-agent-swarm-data-stack/>